

The contact-edge selector: why $\theta_c = \inf \mathcal{F}$ fails, and the closed-form tangency residual that replaces it

SpectralKM derivations

1 The symptom — retracted

An earlier version of this note opened by asserting that the model overpredicts the contact time by roughly a factor of two. That is retracted: it was an artefact of the comparison, not a property of the model.

Run at *identical* dimensionless parameters, this model and the reference implementation (harrislab-brown/BouncingDroplets, executed directly in MATLAB) agree:

metric	ours	reference	difference
contact duration (<code>primary_contact_time</code> / their $f > 0$)	3.6293	3.6530	0.65%
threshold, centre of mass below R	4.4432	4.5297	1.9%
$\min z_{cm}/R$	0.1670	0.1110	we penetrate less

at $We = 1.0959$, $Bo = 0.01661$, $Oh = 0.00617$ ($R = 0.35$ mm water, $U = 0.4759$ m/s).

Where the false symptom came from. The comparison script selected simulation parameters by *binning on We alone*: it split $\log_{10} We$ into equal intervals over the 931 experimental points and took the median We , Bo and Oh of whatever fell in each. The dataset spans $Oh = 0.0139 \dots 0.7865$ and $Bo = 0.0037 \dots 0.4197$, so a single window mixes fluids differing by $33\text{--}54\times$ in Oh , and the experimental t_c values *within* one window already differ by $1.35\text{--}1.48\times$ — comparable to the “discrepancy” being attributed to the model. The residual table then compared a mixed-parameter simulation curve against a median over that mixture. The confound was noted when the script was written and not acted upon for several rounds of investigation.

A related trap, recorded so it is not repeated: the spreadsheet’s “Figure 5” experimental block sits at $Oh \approx 0.042$, $Bo \approx 0.005$, matching *neither* the paper’s water case (0.017, 0.006) nor its oil case (0.056, 0.058), and its `km_model` column is at a fixed $Bo = 0.0189$, $Oh = 0.0304$. The spreadsheet’s figure numbering does not correspond to the manuscript in **references/**, so which experimental subset pairs with which model curve is still unresolved.

What was eliminated along the way. Each by measurement, and each still worth knowing:

Candidate	Test	Outcome
metric definition	<code>threshold_contact_time</code> vs. <code>contact_time</code>	real, and it does differ
drop viscous closure	<code>:reid</code> vs. <code>:lamb</code>	4th digit
bath radius (reflected waves)	<code>debug_bath_radius_sensitivity.jl</code> $b = 6, 12, 20, 30$ with $M \propto b$	flat: 4.9402, 4.9613, 4.9614, 4.9614
contact-edge rule	<code>:crossing</code> (the reference's own rule)	unchanged to 3 digits

$b = 6$ remains unvalidated per `provenance.tex`; it is now measured as *not* affecting the contact time.

2 What is actually wrong

Tracing a contact (`debug_contact_exit_trace.jl`, `debug_feasibility_collapse.jl`):

- the net force $f \leq 0$ on **0 of 221** steps – it decays asymptotically to 1.4×10^{-8} from above and never reverses, so the $f \leq 0$ exit test never fires and contact terminates only at the θ_c floor;
- 63% of the contact duration occurs *after* the droplet has begun rising;
- at $t = 0.009$ the non-intersection predicate is false for all $\theta < 0.0492$ and $\inf \mathcal{F} = 0.0492$; one step later, at $t = 0.010$, it is true even at $\theta = 10^{-4}$, $\inf \mathcal{F}$ collapses by a factor of 200, and the patch never recovers.

None of this explains a contact-time error, because there is none to explain. What follows stands on its own: the rule is degenerate, that degeneracy is real and reproducible, and it produces a physically indefensible patch at low We . It is a defect worth fixing for its own sake, not the cause of a symptom that turned out to be an artefact.

The search is not at fault. A brute-force scan (`debug_selector_rate_limit.jl`) shows `select_theta_c` returning within 0.3–4% of the true $\inf \mathcal{F}$ at every step. The *rule* is at fault.

3 The degeneracy, stated precisely

Write the gap as $C(\theta, \tau) = \eta(r(\theta, \tau), \tau) - z_{cm} + z_d(\theta, \tau)$. The kinematic match imposes $C = 0$ on $[0, \theta_c]$; physical consistency requires $C < 0$ beyond the edge. Continuing just outside, $C(\theta) \simeq a(\theta - \theta_c)$ with

$$a(\theta_c) := \partial_\theta C|_{\theta_c}, \quad (1)$$

so non-overlap is the inequality $a(\theta_c) \leq 0$ and the feasible set is $\mathcal{F} = \{\theta_c : a(\theta_c) \leq 0\}$. The two rules are then

$$\text{current: } \theta_c = \inf \mathcal{F}, \quad \text{tangency (3): } a(\theta_c) = 0. \quad (2)$$

If $a(\theta_c) < 0$ for every $\theta_c > 0$ – the measured situation from $t = 0.010$ onward – then $\mathcal{F} = (0, \theta_{\max}]$ and $\inf \mathcal{F} = 0$, *regardless of the shape or magnitude of a .*

This is the whole failure. The infimum rule uses only the **sign** of a and never its magnitude, so it returns the same answer whether the surfaces are barely separated or widely so. It carries information only while the constraint is *binding*; the moment non-intersection goes inactive, it degenerates to $\theta_c = 0$ – not because the physics says the patch vanishes, but because the rule has

nothing left to determine it with. Verified symbolically in `derivations/cas_tangency_residual.py`, which also exhibits an $a < 0$ example where $\inf \mathcal{F} = 0$ while $\arg \min |a|$ is interior.

Two further facts from the scan bear on the rule’s stated justification. Positivity ($p > 0$) holds *at* the collapsed patch, so complementarity – which `select_theta_c`’s docstring cites as its basis – does not forbid it. And the band’s upper edge is set by $f < 0$ and $p < 0$, i.e. by suction, not by `eq:check-monotone-r`. The rule’s original evidence was a single measurement at one step at onset, precisely where the constraint *is* active and the rule does work.

4 The closed-form tangency residual

`eq:tangency-selector` is $T = \partial_\theta C|_{\theta_c}$, currently evaluated in `src/residual.jl` by a centred finite difference with $h = 10^{-6}$.

A claim withdrawn. An earlier draft of this note asserted that differencing a Fourier–Bessel sum at $h = 10^{-6}$ “discards roughly half the available digits” and was therefore unfit to root-find. *That was wrong, and the test written to check it said so:* against a Richardson-extrapolated reference at $h = 10^{-3}$, the existing finite-difference residual agrees to `rto1` = 10^{-8} at $\theta = 0.25, 0.6, 1.1$ (`test/test_selector.jl`). Accuracy is not the problem. The closed form below is still worth having — it removes a step-size parameter, and it differentiates cleanly under `ForwardDiff` where a difference of a difference would not — but the case for it is structural, not numerical, and the fix to the selector does not depend on it. With $\xi(\theta) = 1 + \sum_l \beta_l P_l(\cos \theta)$, $r = \xi \sin \theta$, $z_d = -\xi \cos \theta$ and $\eta(r) = \sum_m a_m J_0(k_m r)$, the chain rule gives

$$T = \eta'(r) [\xi' \sin \theta + \xi \cos \theta] + [\xi \sin \theta - \xi' \cos \theta] \quad (3)$$

with

$$\xi'(\theta) = -\sin \theta \sum_l \beta_l P'_l(\cos \theta), \quad \eta'(r) = -\sum_m a_m k_m J_1(k_m r). \quad (4)$$

(3) is verified against symbolic differentiation of the composed expression (difference exactly zero), and against the existing numerical residual in `test/test_selector.jl`.

5 The pole degeneracy is physical, not algebraic

$T \equiv 0$ at $\theta = 0$ and $\theta = \pi$ for any state – `eq:tangency-degenerate`. Worth recording *how* the CAS established this, because the first attempt was wrong: with η left as an abstract function the CAS correctly **refused**, returning $T(0) = \xi(0) \eta'(0)$, since nothing in the chain rule makes $\eta'(0)$ vanish. It vanishes because of **axisymmetry**, through $J_1(0) = 0$; and ξ' carries a factor $\sin \theta$, so it vanishes at both poles too. Substituting the concrete expressions, $T \rightarrow 0$ at both poles for $l = 2$ and $l = 3$ and arbitrary $\beta_l, a_m, k_m, z_{cm}$.

Two implementation consequences follow. A seed is required at onset, which `onset_theta_c` already supplies from the geometric crossing $C = 0$; and the residual must be root-found *away* from $\theta = 0$, never bracketed from it.

6 The replacement rule

Select θ_c as a root of (3), retaining `eq:check-nonintersect` and `eq:check-monotone-r` as a *feasibility filter* over candidates – which is what `paper-formulation.tex` describes, and what [1] and [2]

do. Because $\eta(r(\theta))$ is a sum of oscillatory J_0 terms, T can change sign more than once on $(0, \pi)$, so a plain bisection on a sign change is not safe; the argmin form of `eq:theta-c-argmin` over admissible candidates is used instead, seeded from the previous step and, at onset, from `onset_theta_c`.

7 Scope

This note establishes that the current rule is degenerate. It does *not* establish that any replacement improves an observable — and the measurement that was supposed to demonstrate that turned out to be measuring a comparison artefact. `:crossing`, which is the reference’s own rule, changes t_c by less than 0.2%. Any future claim here must be made against *parameter-matched* experimental data, matching Bo and Oh and not We alone. Nor does it revisit the upper edge of the admissible band, where $f < 0$ and $p < 0$ already exclude wide patches; whether that exclusion is itself too aggressive is a separate question.

References

- [1] L. F. L. Alventosa, R. Cimpanu, D. M. Harris, *Inertio-capillary rebound of a droplet impacting a fluid bath*, JFM (2023).
- [2] C. A. Galeano-Ríos, P. A. Milewski, J.-M. Vanden-Broeck, *Non-wetting impact of a sphere onto a bath*, JFM (2017).